



Accelerating query processing with materialized views in Apache Hive

Jesús Camacho Rodríguez

DataWorks Summit San Jose
June 19, 2018

Apache Hive

- Initial use case: batch processing
 - Read-only data
 - HiveQL (SQL-like query language)
 - MapReduce
- Effort to take Hive beyond its batch processing roots
 - Started in Apache Hive 0.10.0 (January 2013)
 - Latest release: Apache Hive 3.0 (May 2018)
- Extensive renovation to improve **three different axes**
 - Latency: allow interactive and sub-second queries
 - Scalability: from TB to PB of data
 - SQL support: move from HiveQL to SQL standard



Apache Hive

Important internals improvements

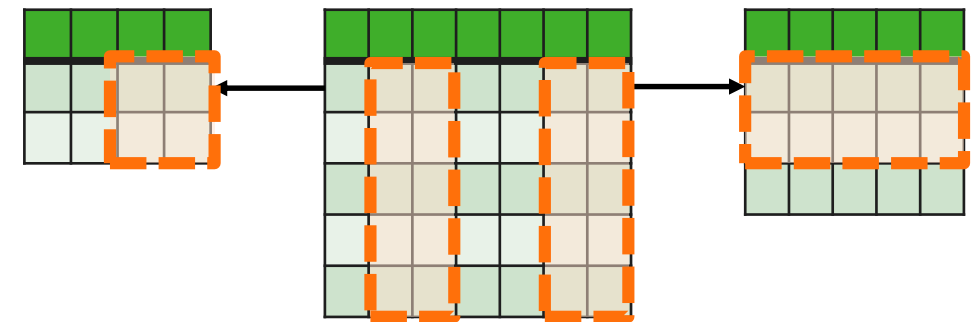
- Multiple execution engines: Apache Tez and Apache Spark
- More efficient join execution algorithms
- Vectorized query execution
 - Integration with columnar storage formats: Apache ORC, Apache Parquet
- LLAP (*Live Long and Process*)
 - Persistent daemons for low-latency queries
- Rule-based and cost-based optimizer
 - Better statistics
- Tighter integration with other data processing systems: Druid



Accelerating query processing

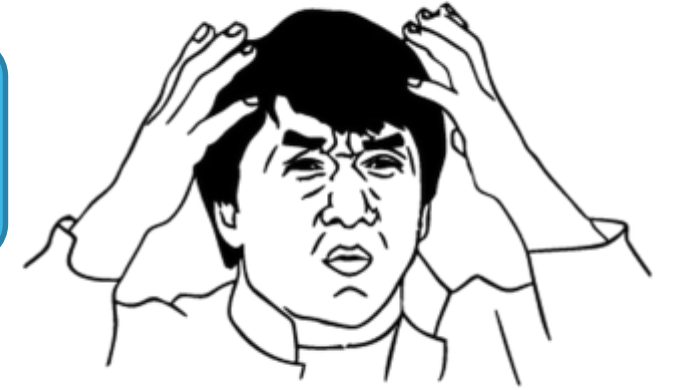
Optimization based on access patterns

- Change data physical properties (distribute, sort)
- Filter rows
- Denormalize
- Preaggregate



Accelerating query processing

Currently, Hive users have to do it manually



Optimization based on access patterns

- Establish relationship between original and new tables
 - Has a similar table already been created?
- Rewrite your queries to use new tables
 - What happens when access patterns change?
- Maintain your new tables when original tables change
 - Do I have to fully rebuild new tables?



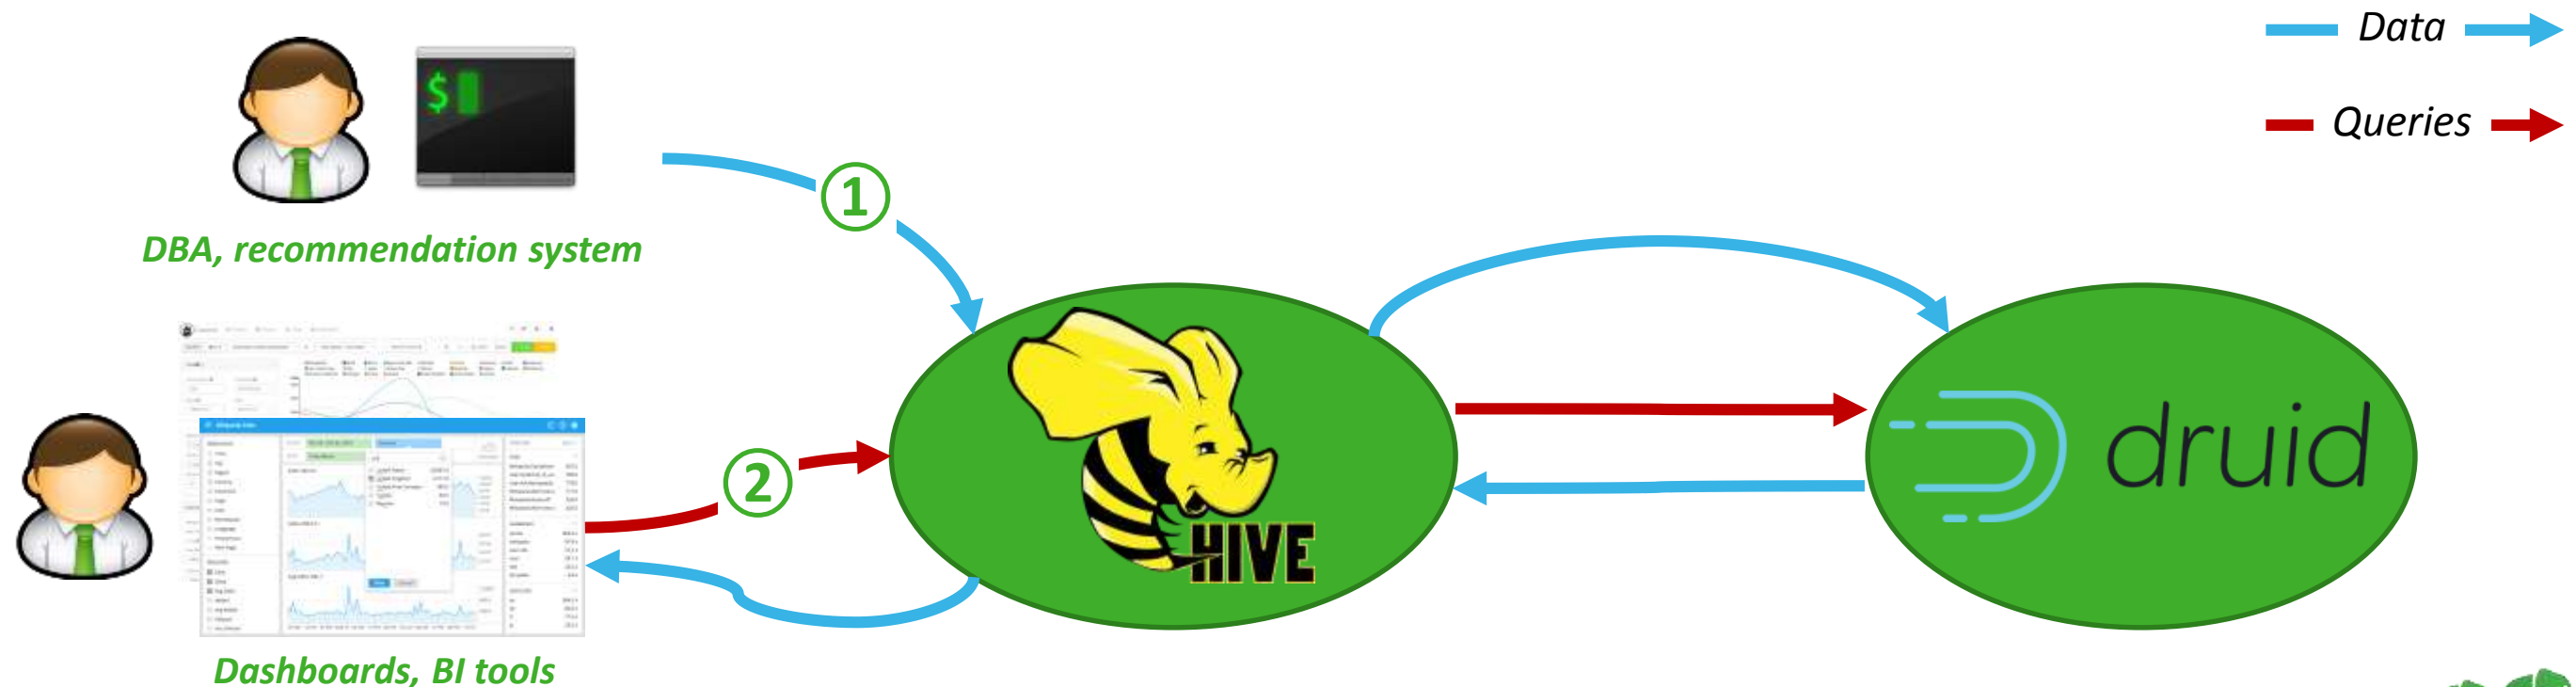
Materialized views

- A **materialized view** is an entity that contains the result of a evaluating a query
 - Important property → Awareness of the materialized view definition semantics
 - Optimizer can exploit them for automatic query rewriting
 - System can handle maintenance of the materialized views
- Generally, materializations can be created in different forms depending on the scope
 - DBA writes “CREATE MATERIALIZED VIEW” statement
 - Daemon creates materialized view based on recent query activity
 - Cached result of previous similar query
 - Query factorization identifies common pieces within a single query

Possible workflow

```
CREATE MATERIALIZED VIEW `ssb_mv`  
STORED AS 'org.apache.hadoop.hive.druid.DruidStorageHandler'  
ENABLE REWRITE  
AS  
<query>;
```

1. Create materialized view using Hive tables
 - Stored by Hive or Druid
2. User or dashboard sends queries to Hive
 - Hive rewrites queries using available materialized views
 - Execute rewritten query



Materialized views in Apache Hive

- First implementation is part of **Apache Hive 3.0**
- Multiple storage options: Hive, Druid
- Automatic rewriting of incoming queries to use materialized views
- Efficient view maintenance
 - Incremental refresh
- Multiple options to control materialized views lifecycle

Management of materialized views in Hive

Materialized view creation

- *CREATE MATERIALIZED VIEW* statement

```
CREATE MATERIALIZED VIEW [IF NOT EXISTS] [db_name.]materialized_view_name
[DISABLE REWRITE]
[COMMENT materialized_view_comment]
[
  [ROW FORMAT row_format]
  [STORED AS file_format]
  | STORED BY 'storage.handler.class.name' [WITH SERDEPROPERTIES (...)]
]
[LOCATION hdfs_path]
[TBLPROPERTIES (property_name=property_value, ...)]
AS
<query>;
```

---> Supports custom table properties, storage format, etc.



Materialized view creation (stored in Druid)

- *CREATE MATERIALIZED VIEW* statement

```
CREATE MATERIALIZED VIEW druid_wiki_mv
STORED AS 'org.apache.hadoop.hive.druid.DruidStorageHandler'
AS
SELECT __time, page, user, c_added, c_removed
FROM src;
```

Hive materialized view name

Hive storage handler classname



Other operations for materialized view management

DROP MATERIALIZED VIEW [db_name.]materialized_view_name;

SHOW MATERIALIZED VIEWS [IN database_name] ['identifier_with_wildcards'];

DESCRIBE [EXTENDED | FORMATTED] [db_name.]materialized_view_name;

→ More operations to be added and extended

Materialized view-based query rewriting

Materialized view-based rewriting algorithm



- Automatically rewrite incoming queries using materialized views
 - Optimizer exploits materialized view definition semantics
- Built on the ideas presented in [GL01] using **Apache Calcite**
 - Supports queries containing TableScan, Project, Filter, Join, Aggregate operators
- Includes some extensions
 - Generation of additional rewritings without needing to do join permutation
 - Partial rewritings using union operators
- More information about the rewriting coverage
 - http://calcite.apache.org/docs/materialized_views#rewriting-using-plan-structural-information

[GL01] Jonathan Goldstein and Per-åke Larson. Optimizing queries using materialized views: A practical, scalable solution. In Proc. ACM SIGMOD Conf., 2001.

Enable materialized view-based rewriting

- Global property to enable materialized view rewriting for queries

```
SET hive.materializedview.rewriting=true;
```

- User can selectively use enable/disable materialized views for rewriting
- Materialized views are enabled by default for rewriting
- Behavior can be altered after materialized view has been created

```
ALTER MATERIALIZED VIEW [db_name.]materialized_view_name ENABLE|DISABLE REWRITE;
```

Materialized view-based rewriting (example)

- Materialized view definition

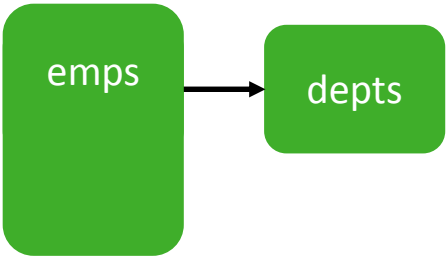
Employees that were hired after 2016

```
CREATE MATERIALIZED VIEW mv
AS
SELECT empid, deptname, hire_date
FROM emps JOIN depts
ON (emps.deptno = depts.deptno)
WHERE hire_date >= '2016-01-01';
```

- Query

Employees that were hired last quarter

```
SELECT empid, deptname
FROM emps JOIN depts
ON (emps.deptno = depts.deptno)
WHERE hire_date >= '2018-01-01'
AND hire_date <= '2018-03-31';
```



- Materialized view-based rewriting

```
SELECT empid, deptname
FROM mv
WHERE hire_date >= '2018-01-01'
AND hire_date <= '2018-03-31';
```

mv contents

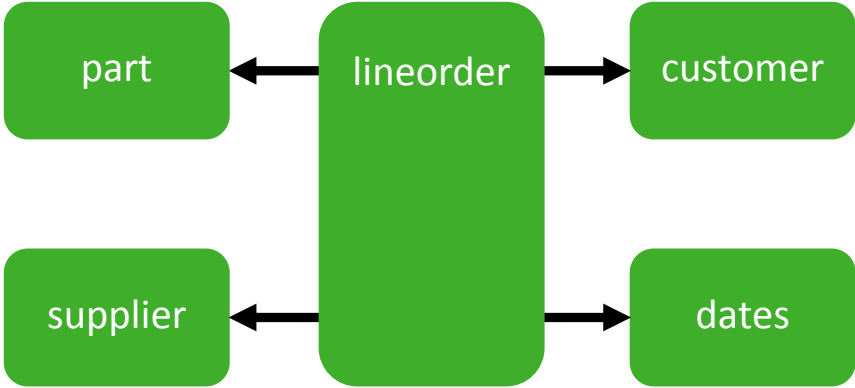
empid	deptname	hire_date
10001	IT	2016-03-01
10002	IT	2017-01-02
10003	HR	2017-07-01
10004	Finance	2018-01-15
10005	HR	2018-02-02

Query results

empid	deptname
10004	Finance
10005	HR



Materialized view-based rewriting (example 2)



- Materialized view definition

```
CREATE MATERIALIZED VIEW mv AS
SELECT <dims>,
       lo_revenue,
       lo_extendedprice * lo_discount AS d_price,
       lo_revenue - lo_supplycost
FROM   customer, dates, lineorder, part, supplier
WHERE  lo_orderdate = d_datekey
       and lo_partkey = p_partkey
       and lo_suppkey = s_suppkey
       and lo_custkey = c_custkey;
```

- Query

```
SELECT sum(lo_extendedprice * lo_discount)
FROM   lineorder, dates
WHERE  lo_orderdate = d_datekey
       and d_year = 2013
       and lo_discount between 1 and 3;
```

- Materialized view-based rewriting

```
SELECT SUM(d_price)
FROM   mv
WHERE  d_year = 2013
       and lo_discount between 1 and 3;
```

Exploit SQL PK-FK and NOT NULL constraints

mv contents

d_year	lo_discount	<dims>	d_price
2013	2	...	7.55
2014	4	...	432.60
2013	2	...	34.45
2012	2	...	2.05
...

Query results



sum
42.0
...



Materialized view-based rewriting (example 3)



- Materialized view definition

```
CREATE MATERIALIZED VIEW mv AS
SELECT floor(time to minute), page,
       SUM(added) AS c_added,
       SUM(removed) AS c_rmv
FROM wiki
GROUP BY floor(time to minute), page;
```

- Query

```
SELECT floor(time to month),
       SUM(added) AS c_added
FROM wiki
GROUP BY floor(time to month);
```

- Materialized view-based rewriting

```
SELECT floor(time to month),
       SUM(c_added) as c_added
FROM mv
GROUP BY floor(time to month);
```

mv contents

__time	page	c_added	c_rmv
2011-01-01 01:05:00	Justin	1800	25
2011-01-20 19:00:00	Justin	2912	42
2011-01-01 11:06:00	Ke\$ha	1953	17
2011-02-02 13:15:00	Ke\$ha	3194	170
2011-01-02 18:00:00	Miley	2232	34



Query results

__time	c_added
2011-01-01 00:00:00	8897
2011-02-01 00:00:00	3194



Materialized view maintenance

Rebuilding materialized views

- Rebuild needs to be triggered manually by user
`ALTER MATERIALIZED VIEW [db_name.]materialized_view_name REBUILD;`
- **Incremental** materialized view maintenance
 - Only refresh data that has changed in source tables
 - Multiple benefits
 - Decrease rebuild step execution time
 - Preserves LLAP cache for existing data
 - Materialized view should only use transactional tables (micromanaged or ACID)
 - Current implementation only supports incremental rebuild for insert operations
 - Update/delete operations force full rebuild
- Optimizer will attempt incremental rebuild
 - Otherwise, fallback to **full rebuild** (`INSERT OVERWRITE` with MV definition)

Incremental view maintenance algorithm

- Relies on materialized view rewriting algorithm
 - Materialized view stores write ID for its tables when it is created/refreshed
 - Write ID associates rows with transactions
 - When rebuild is triggered, introduce filter condition on write ID column in MV definition
 - Read only new rows from source tables
 - Execute materialized view rewriting
 - Rewrite `INSERT OVERWRITE` (full rebuild) into more efficient plan
 - `INSERT` (table scan, filter, project, join)
 - `MERGE` (table scan, filter, project, join, aggregate)

Incremental view maintenance algorithm (example)

```
CREATE MATERIALIZED VIEW mv1 AS
SELECT page, user,
       SUM(added) AS c_added,
       SUM(removed) AS c_rmv
FROM wiki
GROUP BY page, user;
```

New records



wiki contents

page	user	...	added	removed	...	writeID
...
Miley	Ashu	...	68	16	...	10000
Justin	Zaka	...	392	239	...	10000

```
---> ALTER MATERIALIZED VIEW mv1 REBUILD;
```

mv1 contents

page	user	c_added	c_rmv
Justin	Boxer	1800	25
Justin	Reach	2912	42
Ke\$ha	Xeno	1953	17
Ke\$ha	Helz	3194	170
Miley	Ashu	2232	34



Incremental view maintenance algorithm (example)

```
CREATE MATERIALIZED VIEW mv1 AS
SELECT page, user,
       SUM(added) AS c_added,
       SUM(removed) AS c_rmv
FROM wiki
GROUP BY page, user;
```

New records



wiki contents

page	user	...	added	removed	...	writeID
...
Miley	Ashu	...	68	16	...	10000
Justin	Zaka	...	392	239	...	10000

① Rebuild statement rewriting

```
INSERT OVERWRITE mv1
SELECT page, user, SUM(added) AS c_added, SUM(removed) AS c_rmv
FROM (
  SELECT page, user, c_added, c_removed
  FROM mv1
  UNION ALL
  SELECT page, user, SUM(added) AS c_added, SUM(removed) AS c_rmv
  FROM wiki
  WHERE writeID > 9999
  GROUP BY page, user) subq
GROUP BY page, user;
```

Rollup data

mv1 contents

page	user	c_added	c_rmv
Justin	Boxer	1800	25
Justin	Reach	2912	42
Ke\$ha	Xeno	1953	17
Ke\$ha	Helz	3194	170
Miley	Ashu	2232	34



Incremental view maintenance algorithm (example)

```
CREATE MATERIALIZED VIEW mv1 AS
SELECT page, user,
       SUM(added) AS c_added,
       SUM(removed) AS c_rmv
FROM wiki
GROUP BY page, user;
```

New records →

wiki contents

page	user	...	added	removed	...	writeID
...
Miley	Ashu	...	68	16	...	10000
Justin	Zaka	...	392	239	...	10000

② Rewrite INSERT OVERWRITE into MERGE statement

```
MERGE INTO mv1
USING (
  SELECT page, user, SUM(added) AS c_added, SUM(removed) AS c_rmv
  FROM wiki
  WHERE writeID > 9999
  GROUP BY page, user) src
ON mv1.page = src.page AND mv1.user = src.user
WHEN MATCHED
  THEN UPDATE SET c_added = mv1.c_added + src.c_added,
                  c_removed = mv1.c_removed + src.c_rmv
WHEN NOT MATCHED
  THEN INSERT VALUES (page, user, c_added, c_rmv);
```

mv1 contents

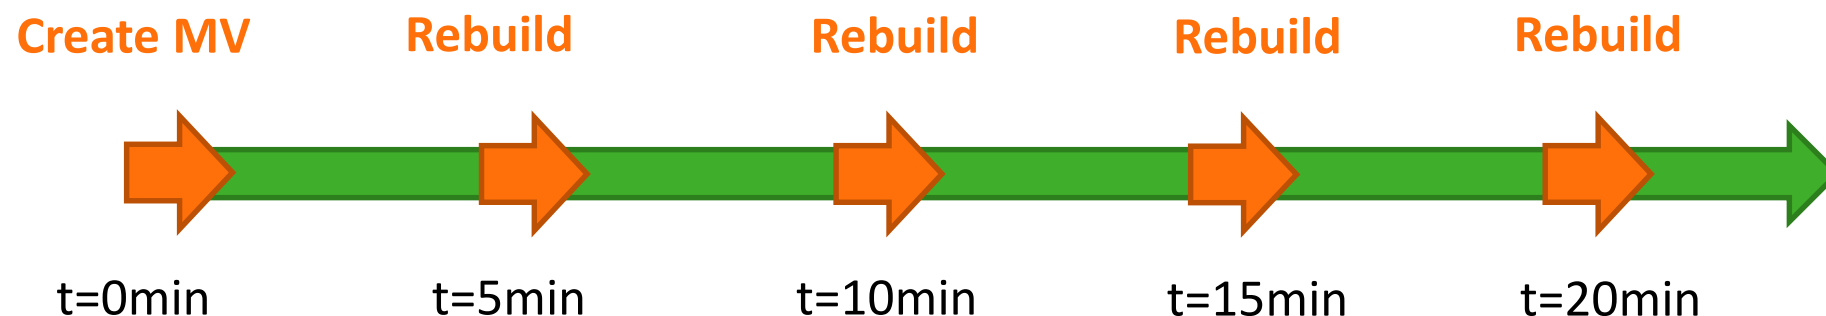
page	user	c_added	c_rmv
Justin	Boxer	1800	25
Justin	Reach	2912	42
Ke\$ha	Xeno	1953	17
Ke\$ha	Helz	3194	170
Miley	Ashu	2300	50
Justin	Zaka	392	239



Materialized view lifecycle

Management of materialized view lifecycle

- Do not accept stale data (default)
 - If content of the materialized view is not fresh, we do not use it for automatic query rewriting
 - Still possible to trigger partial rewritings that read both the stale materialized view and new data from source tables
- Accept stale data
 - *Freshness* defined as a time parameter
 - If MV was not rebuilt for a certain time period and there were changes in base tables, ignore
 - **SET** `hive.materializedview.rewriting.time.window=10min;`
 - Can also be overridden by a certain materialized view using table properties
 - Periodically rebuild materialized view, e.g., every 5 minutes



Road ahead

Road ahead

- Improvements to current materialized views implementation
 - Rewriting performance and scalability
 - Single/many MVs
 - Control physical distribution of data
 - `PARTITIONED BY, DISTRIBUTE BY, SORT BY, CLUSTER BY`
 - Increase incremental view maintenance coverage
 - Support update/delete in source tables
- Materialized view recommender
 - Ease the identification of access patterns for a given workload



Thank you

<https://cwiki.apache.org/confluence/display/Hive/Materialized+views>