LLAP: Sub-Second Analytical Queries in Hive

Gopal Vijayaraghavan

Why LLAP?

- People like Hive
- Disk->Mem is getting further away
  - Cloud Storage isn't co-located
  - Disks are connected to the CPU via network
- Security landscape is changing
  - Cells & Columns are the new security boundary, not files
  - Safely masking columns needs a process boundary
- Concurrency, Performance & Scale are at conflict
  - Concurrency at 100k queries/hour
  - Latencies at 2-5 seconds/query
  - Petabyte scale warehouses (with terabytes of "hot" data)

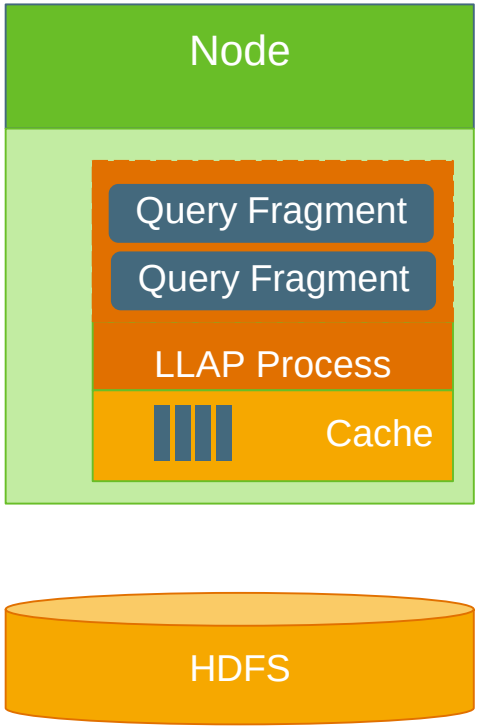

What is LLAP?

- Hybrid model combining daemons and containers for fast, concurrent execution of analytical workloads (e.g. Hive SQL queries)
  - Concurrent queries without specialized YARN queue setup
  - Multi-threaded execution of vectorized operator pipelines
- Asynchronous IO and efficient in-memory caching
- Relational view of the data available thru the API
  - High performance scans, execution code pushdown
  - Centralized data security

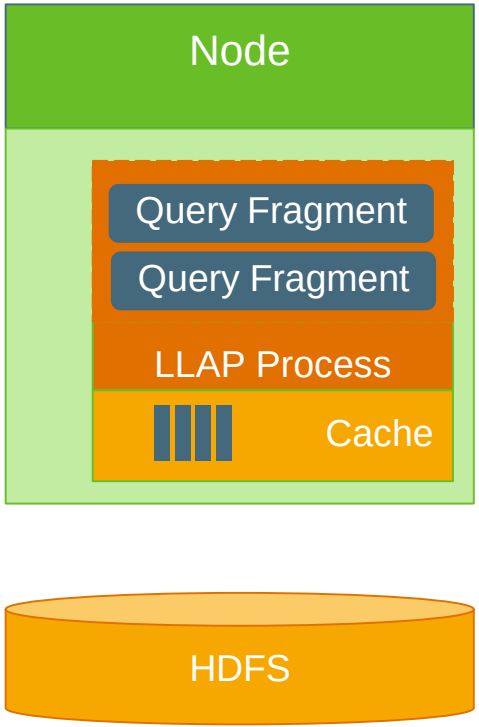

Hortonworks

Hive 2.0 (+ LLAP)

- Transparent to Hive users, BI tools, etc.
- Hive decides where query fragments run (LLAP, Container, AM) based on configuration, data size, format, etc.
- Each Query coordinated independently by a Tez AM
- Number of concurrent queries throttled by number of active AMs
- Hive Operators used for processing
- Tez Runtime used for data transfer

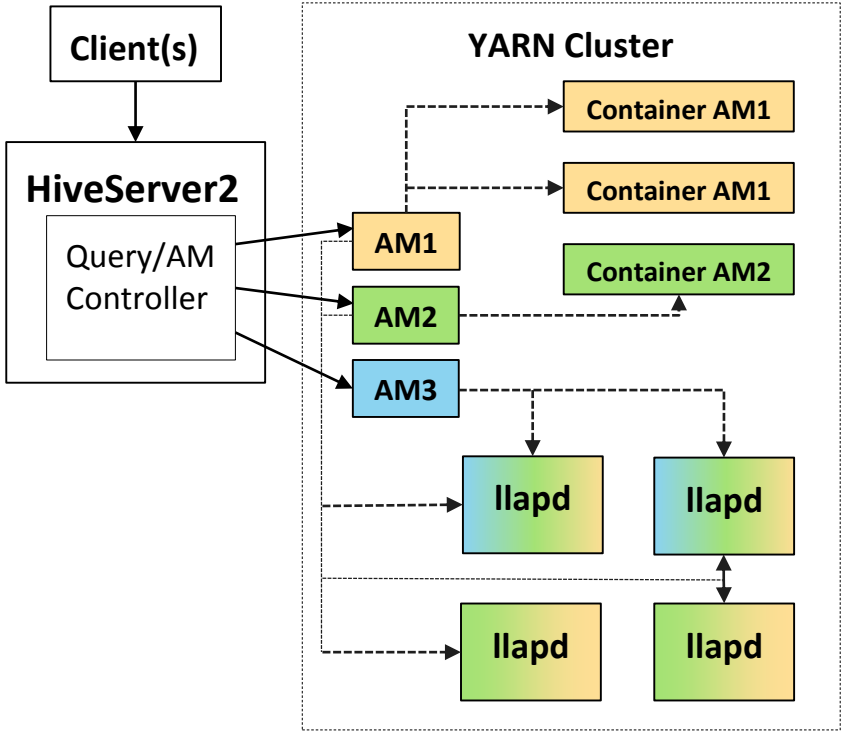

Hortonworks

Industry benchmark – 10Tb scale

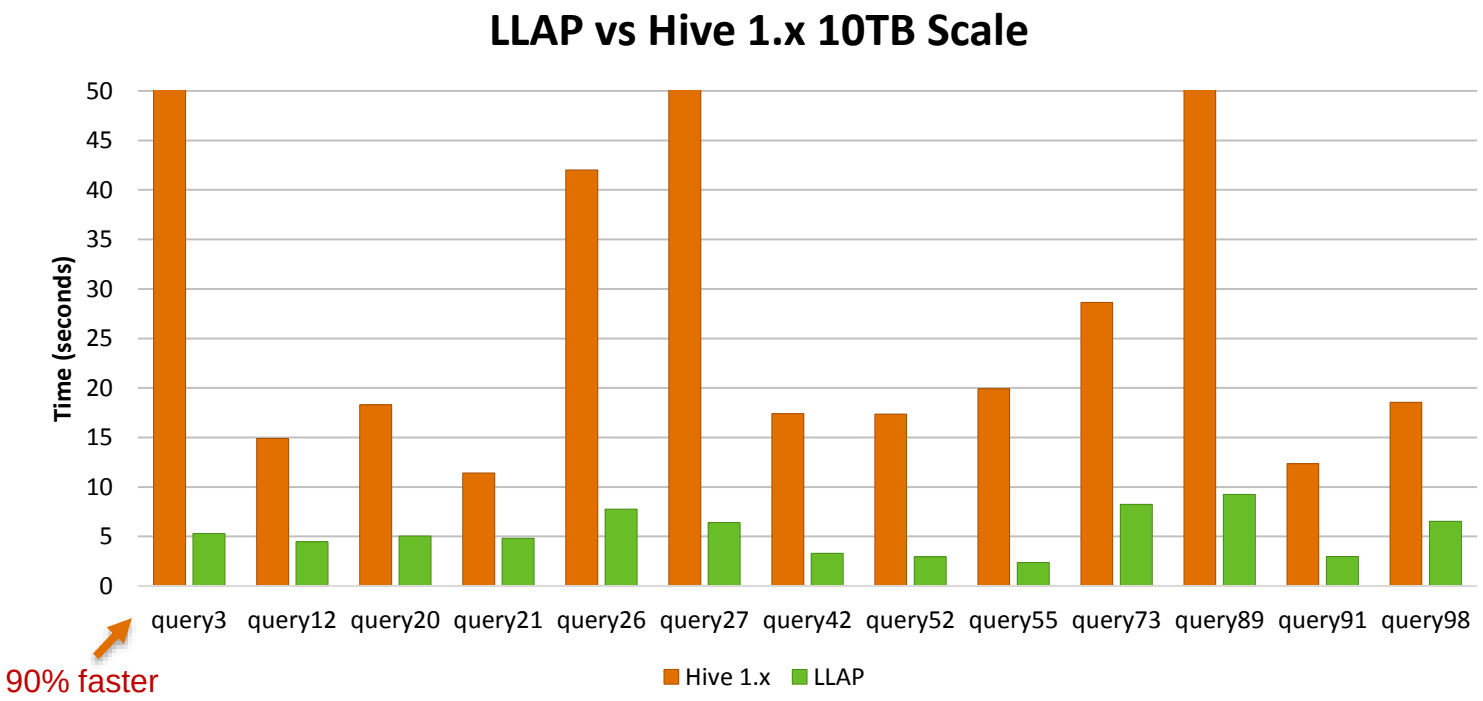

Hortonworks

Evaluation from a customer case study

- Aggregate daily statistics for a time interval:

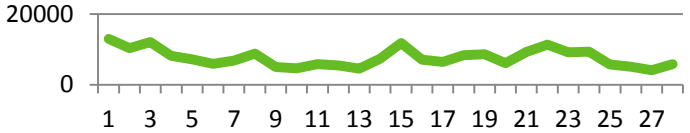

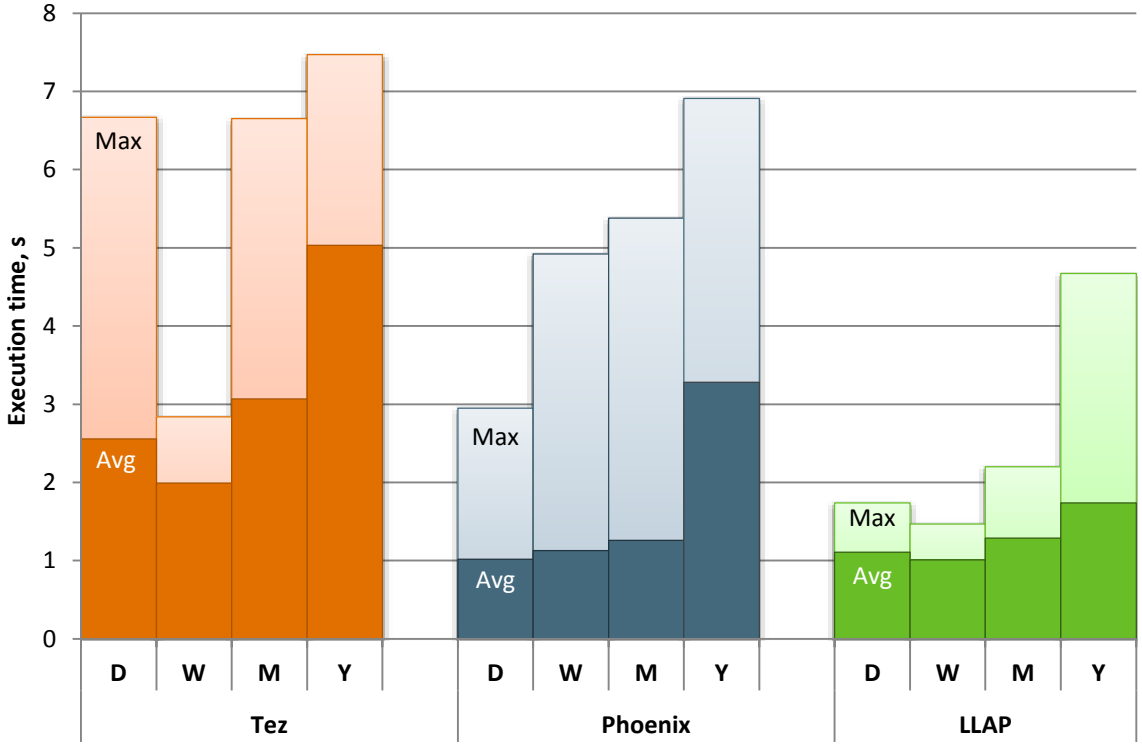


```
SELECT yyyymmdd,
  sum(total_1),
  sum(total_2),
            ...
  from table
  where yyyymmdd >= xxx
            and yyyymmdd < xxx
   and userid = xxx
  group by userid, yyyymmdd;
```


Hortonworks

Evaluation from a customer case study

- Display a large report

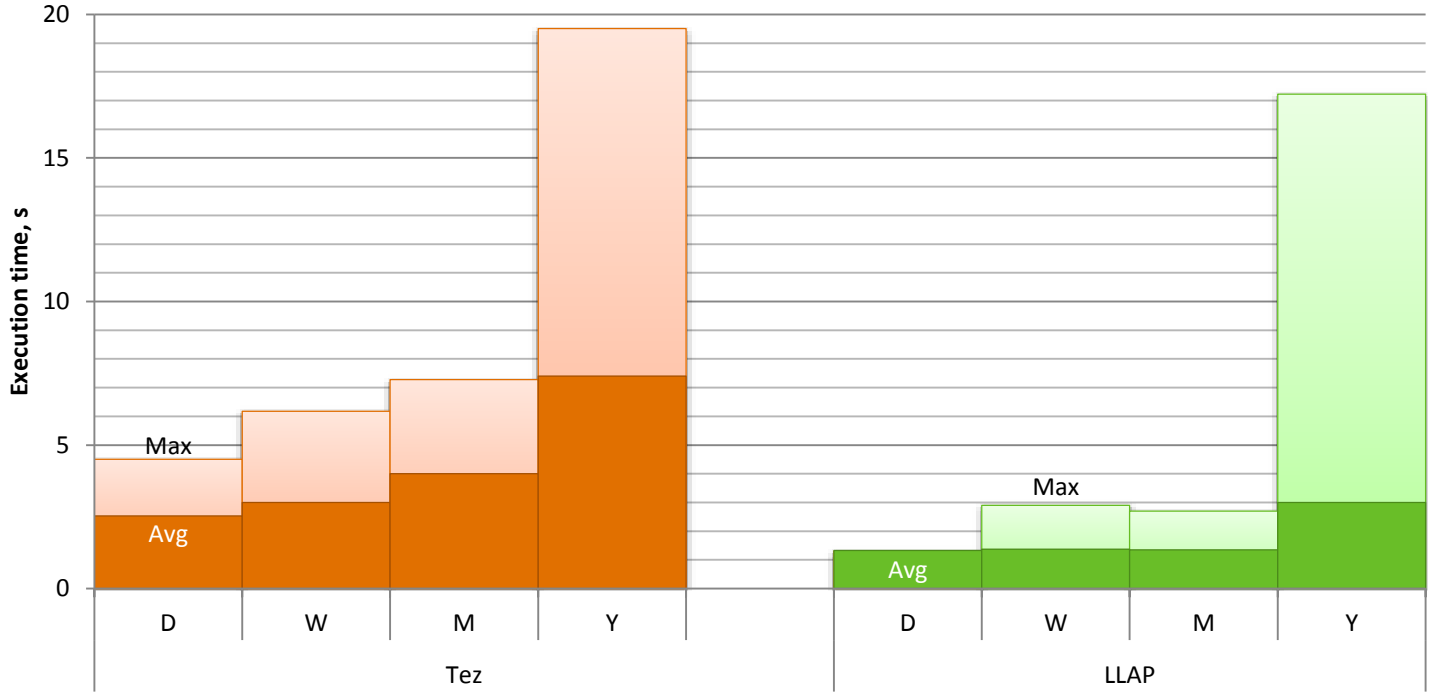

Hortonworks

Cut-away to demo (GIF)

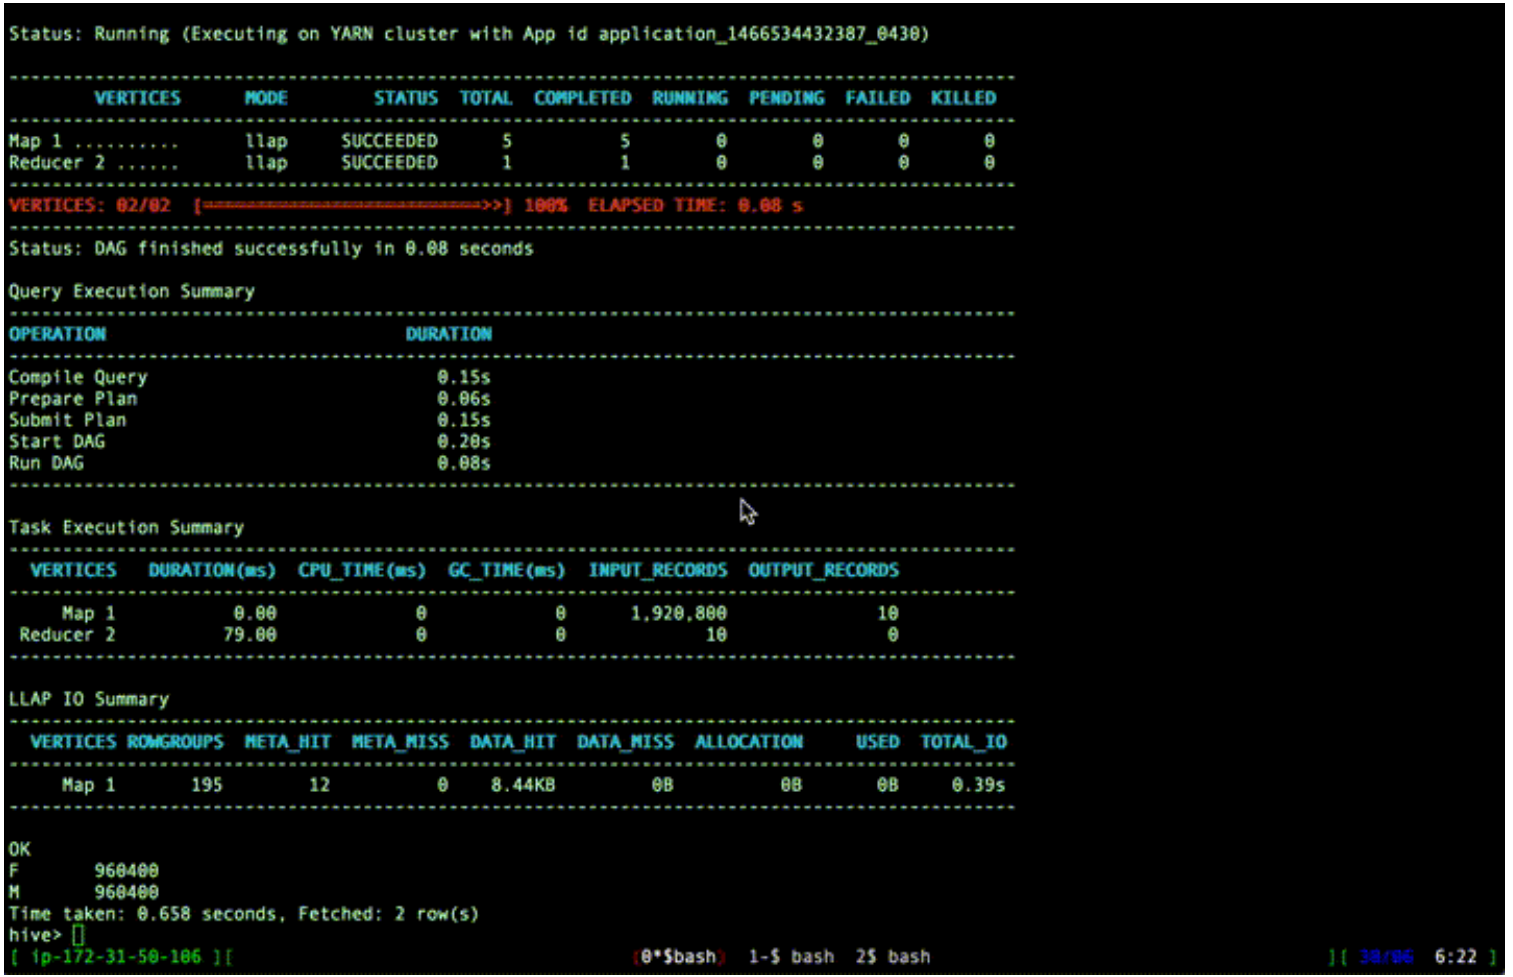

Hortonworks

Hortonworks

How does LLAP make queries faster?

Technical overview – execution

- LLAP daemon has a number of executors (think containers) that execute work "fragments"
- Fragments are parts of one, or multiple parallel workloads (e.g. Hive SQL queries)
- Work queue with pluggable priority
  - Geared towards low latency queries over long-running queries (by default)
- I/O is similar to containers – read/write to HDFS, shuffle, other storages and formats
- Streaming output for data API

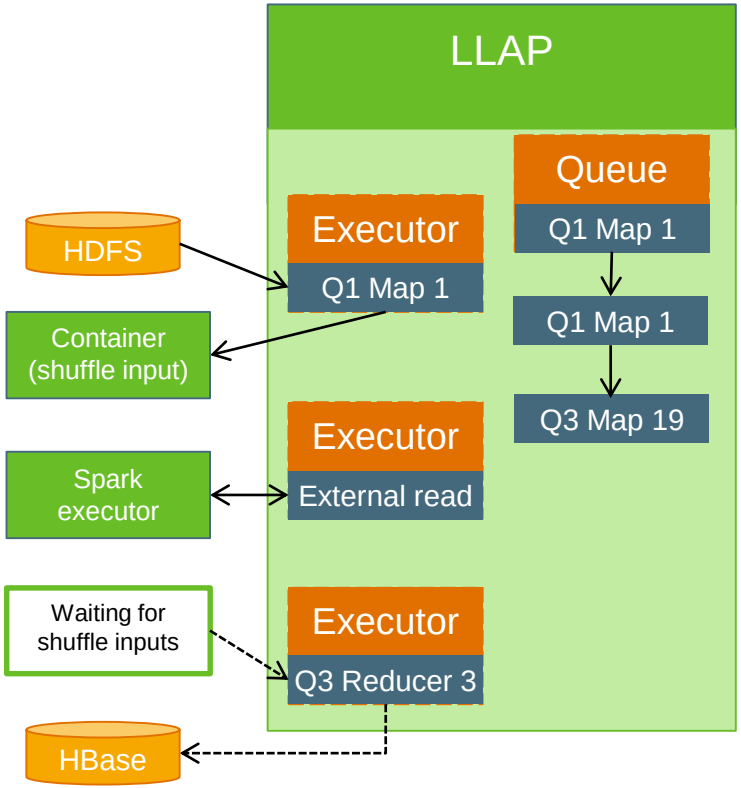

Hortonworks

Technical overview – IO layer

- **Optional:** when executing inside LLAP
  - All other formats use in-sync mode
- Asynchronous IO for Hive
  - Wraps over InputFormat, reads through cache
  - Supported with ORC
- Transparent, compressed in-memory cache
  - Format-specific, extensible
  - NVMe/NVDIMM caches

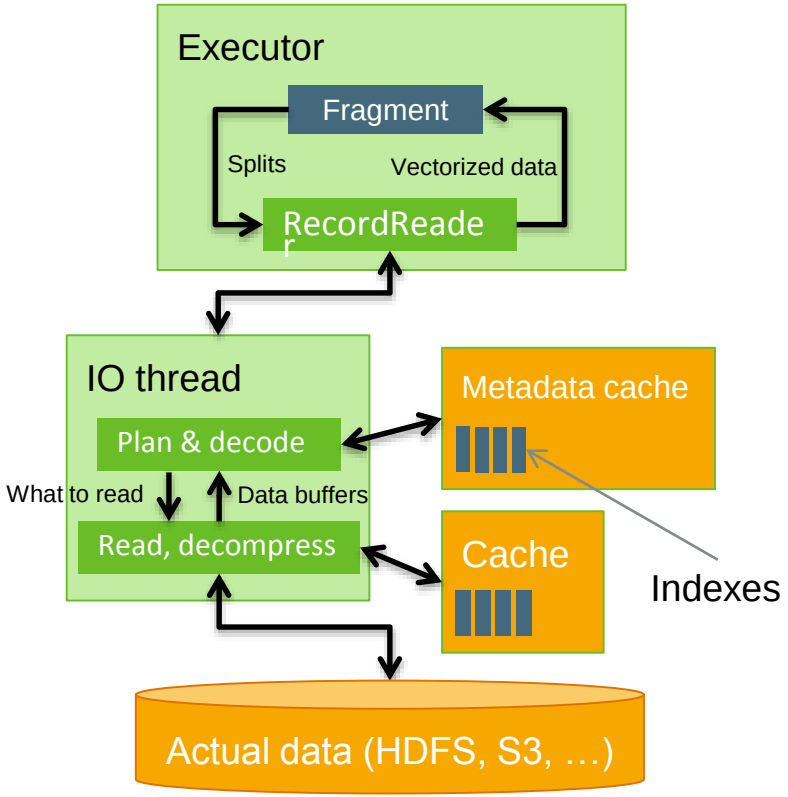

Parallel queries – priorities, preemption

- Lower-priority fragments can be preempted
  - For example, a fragment can start running before its inputs are ready, for better pipelining; such fragments may be preempted
- LLAP work queue examines the DAG parameters to give preference to interactive (BI) queries

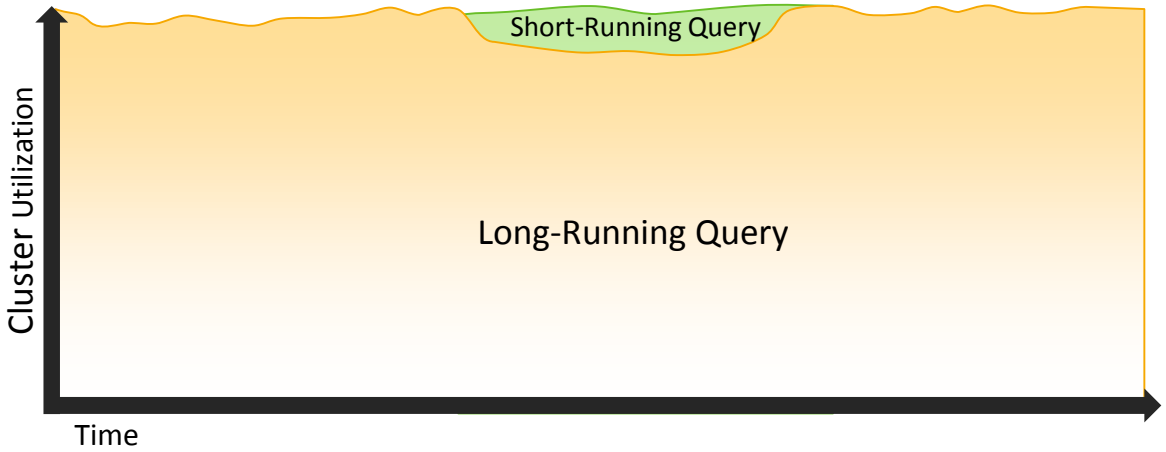

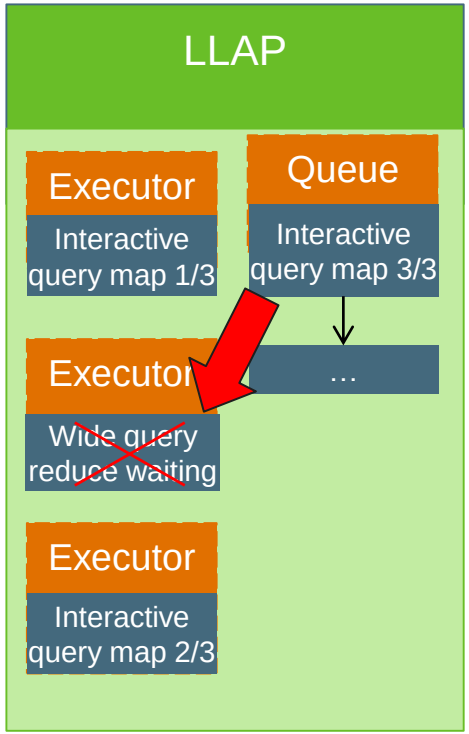

In-memory processing – present and future

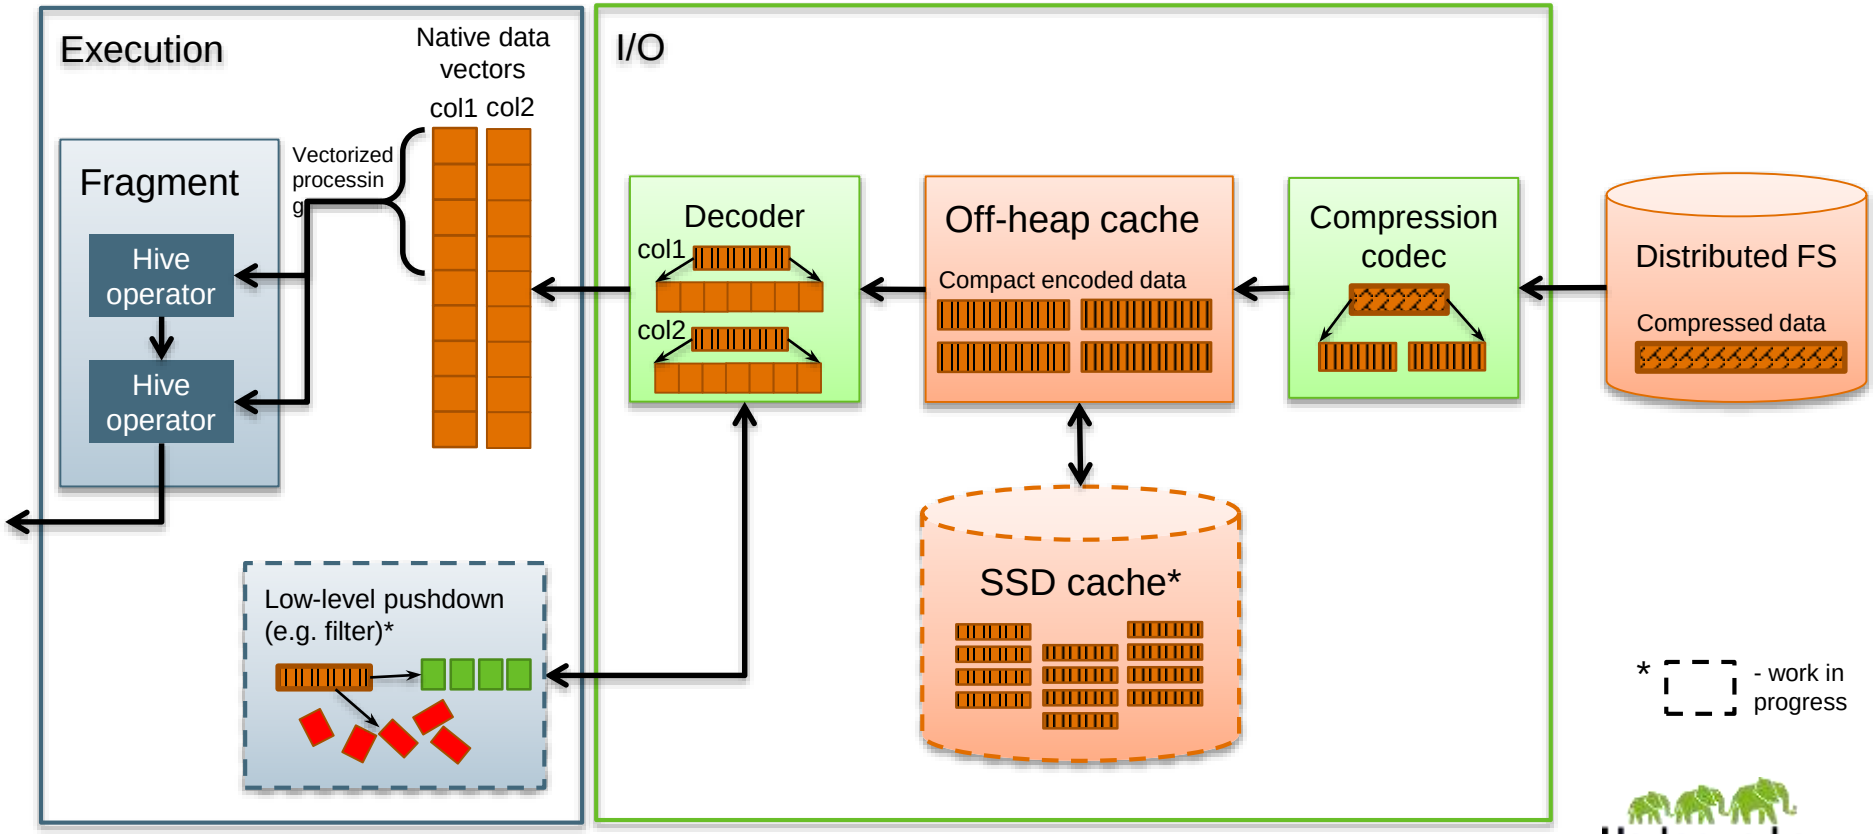

Hortonworks

First query erformance

- Cold LLAP is nearly as fast as shared pre-warmed containers (impractical on real clusters)
- Realistic (long-running) LLAP ~3x faster than realistic (no prewarm) Tez

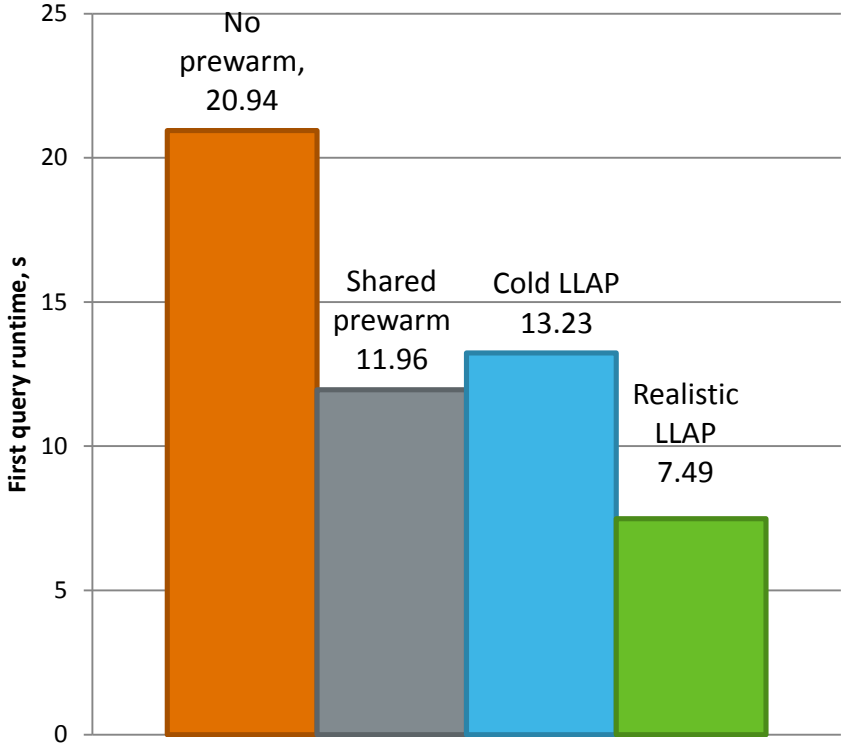

JIT Performance – heavy use

- Cache disabled!

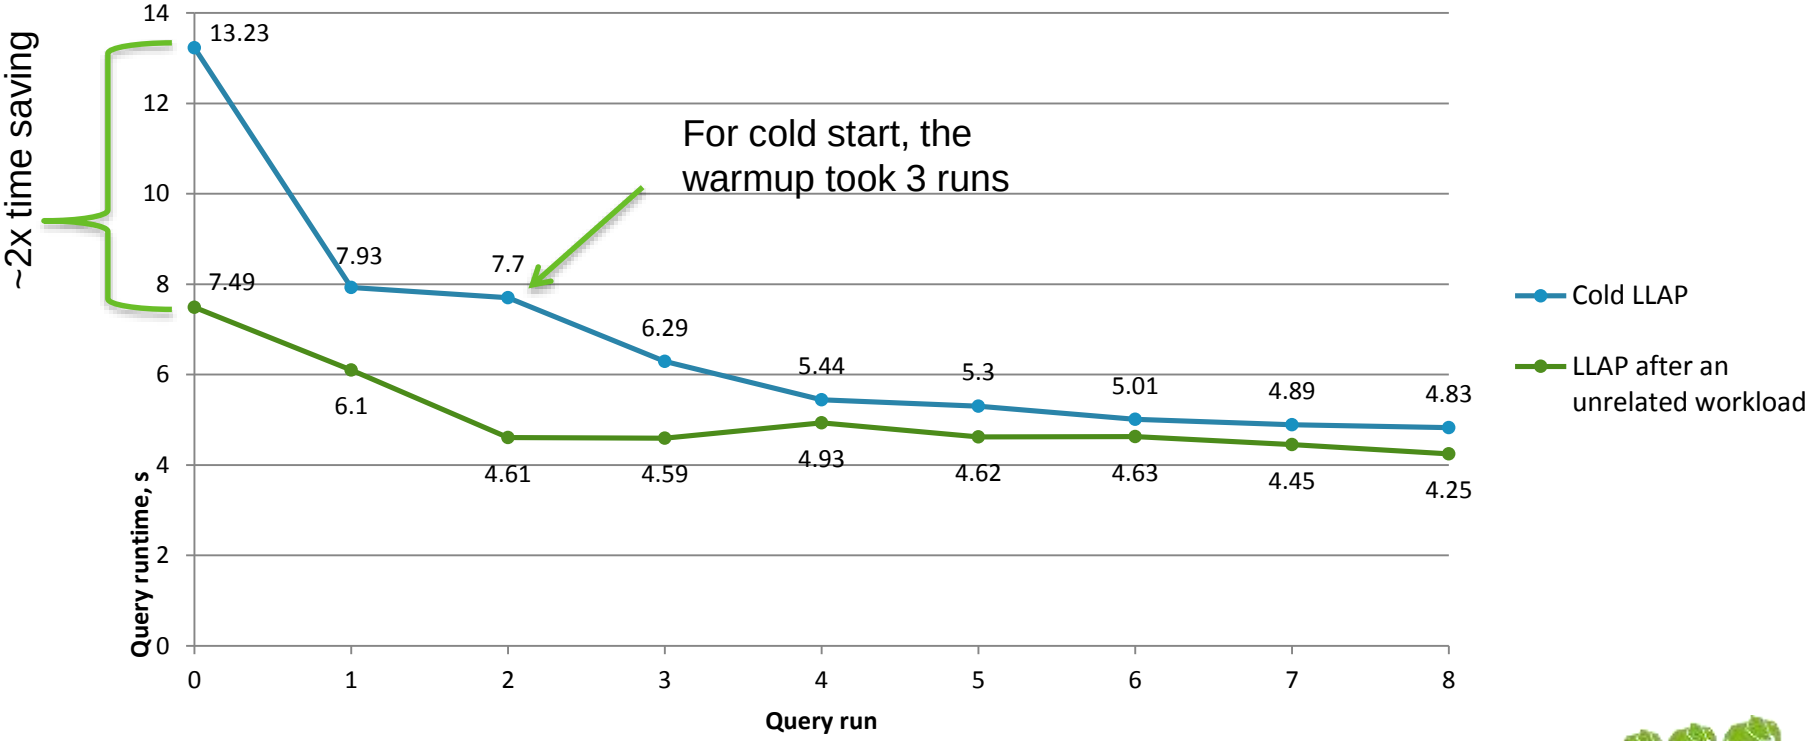

Hortonworks

Parallel query execution – LLAP vs Hive 1.2

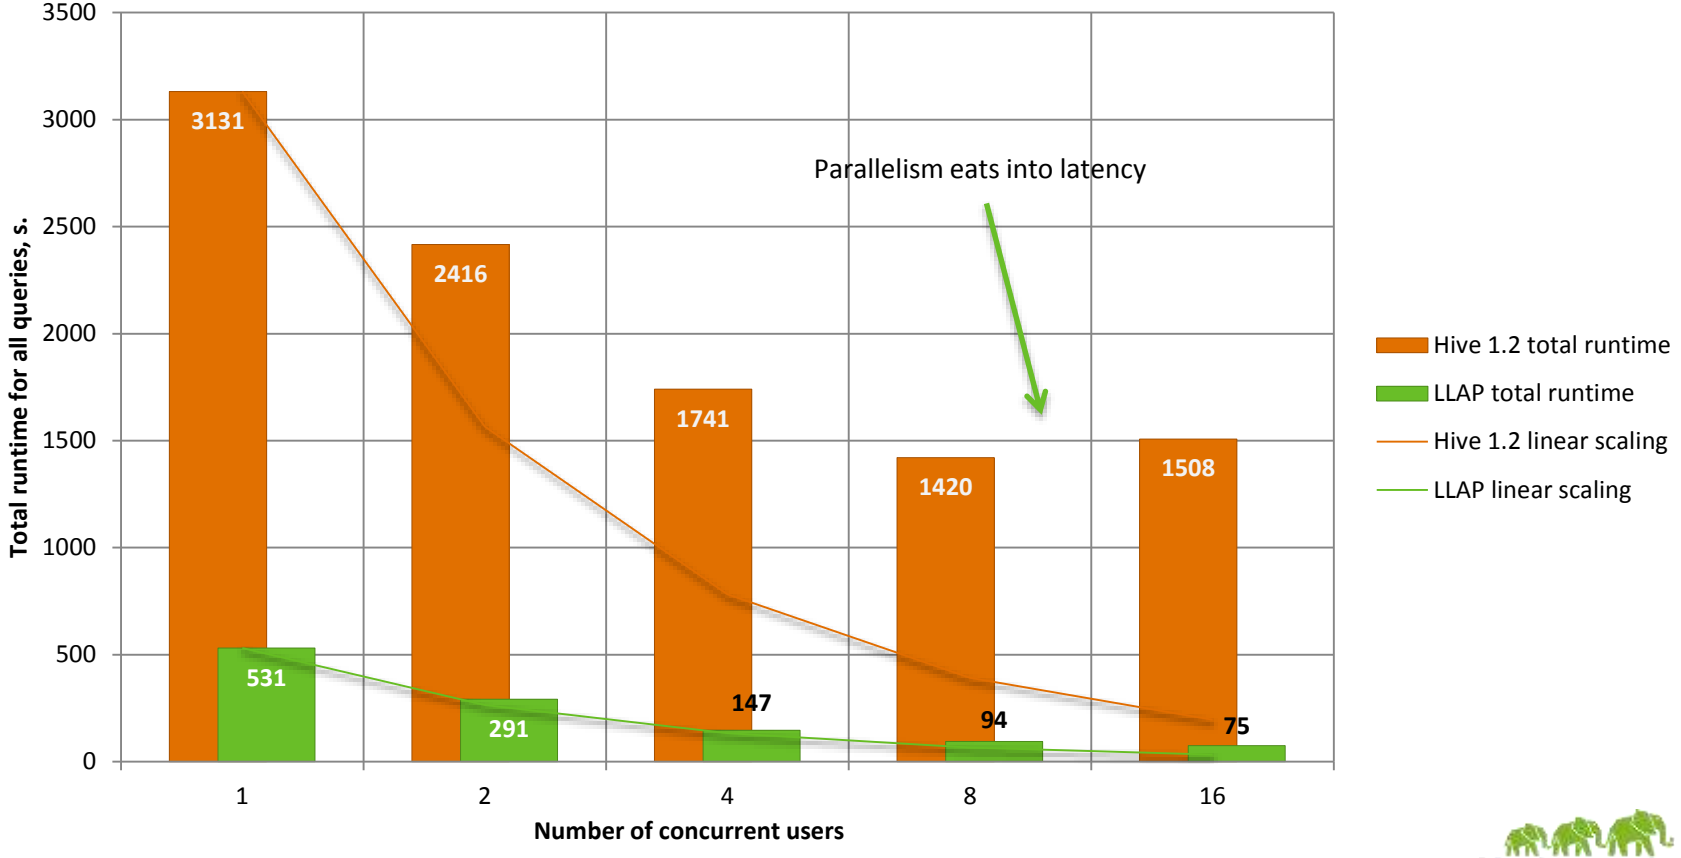

Hortonworks

Parallel query execution – 10Tb scale

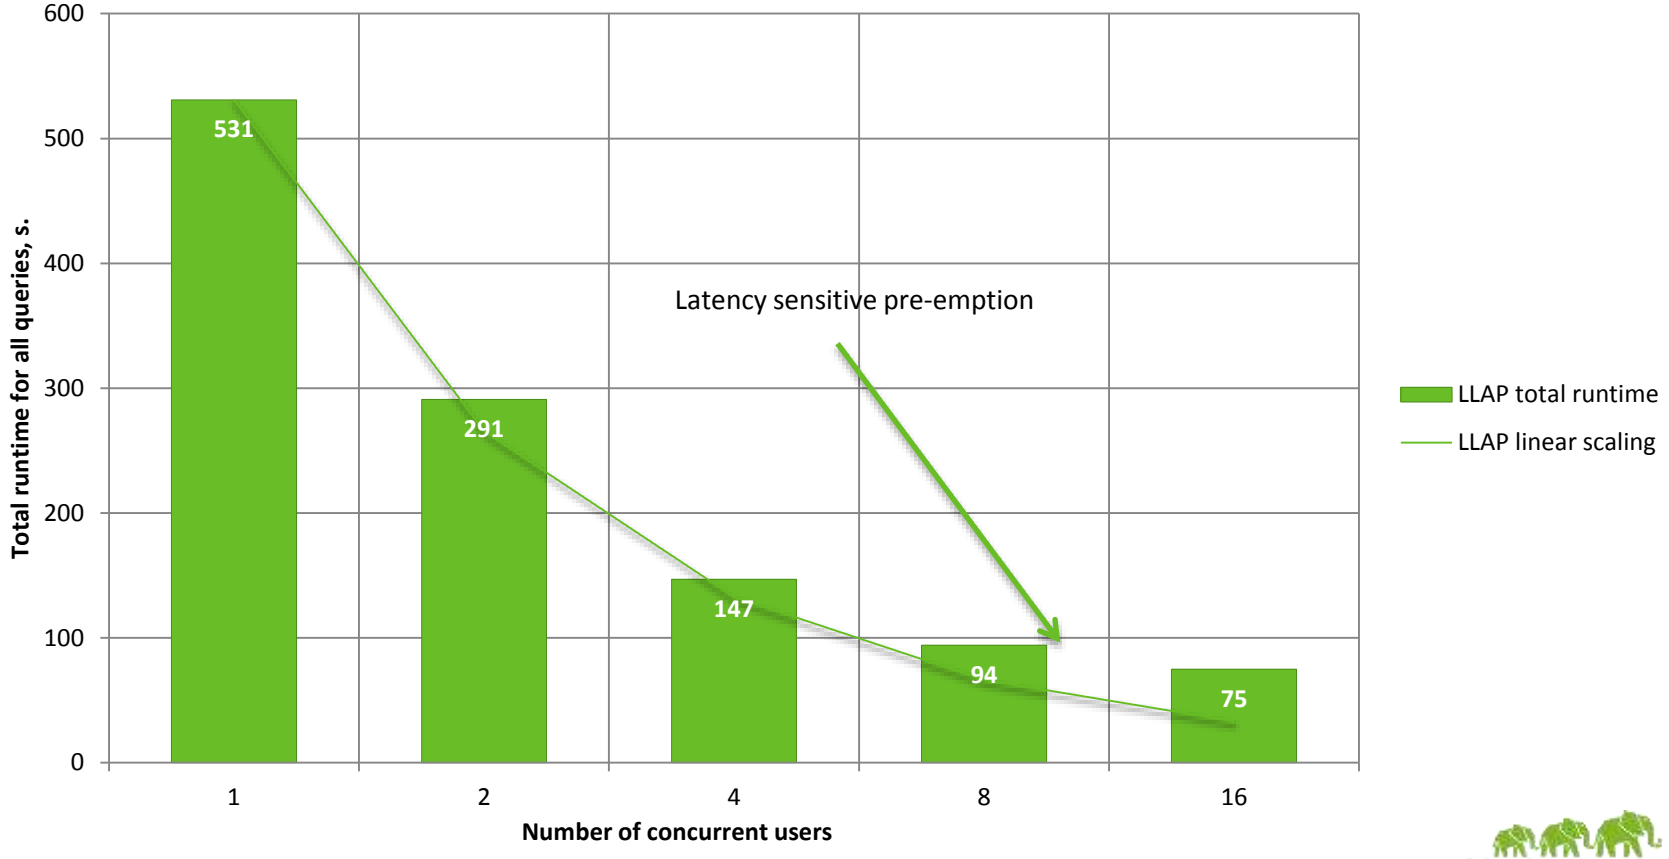

Hortonworks

Performance – cache on HDFS, 1Tb scale

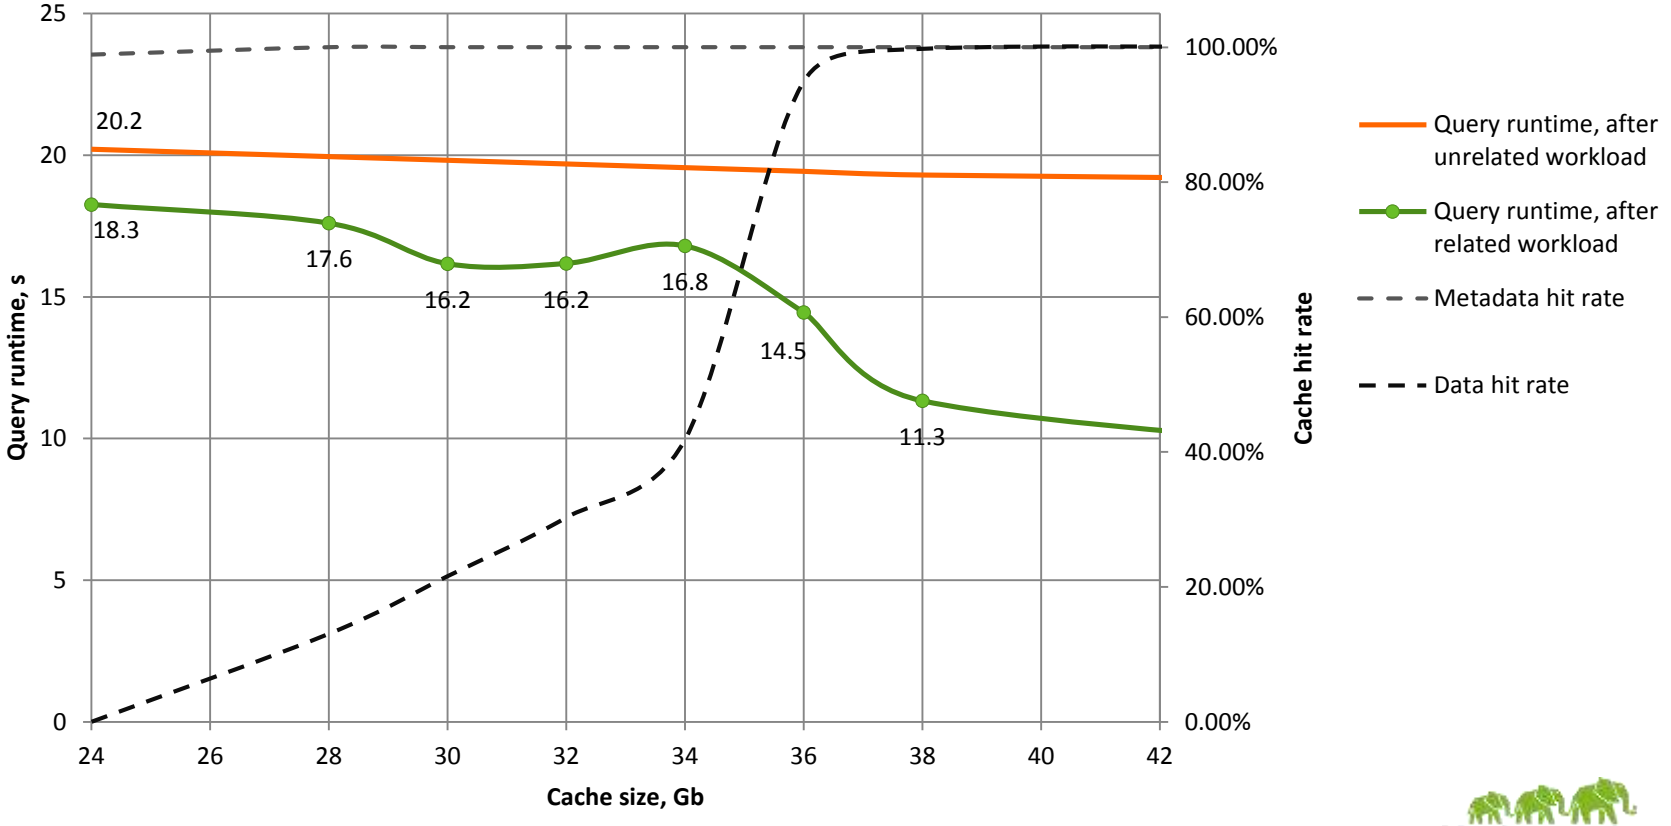

Hortonworks

Hortonworks

LLAP as a “relational” datanode

Example - SparkSQL integration – execution flow

```
var llapContext =
LlapContext.newInstance(
sparkContext, jdbcUrl)

var df: DataFrame =
llapContext.sql("select *
from tpch_text_5.region")
```

DataFrame for Hive/LLAP data

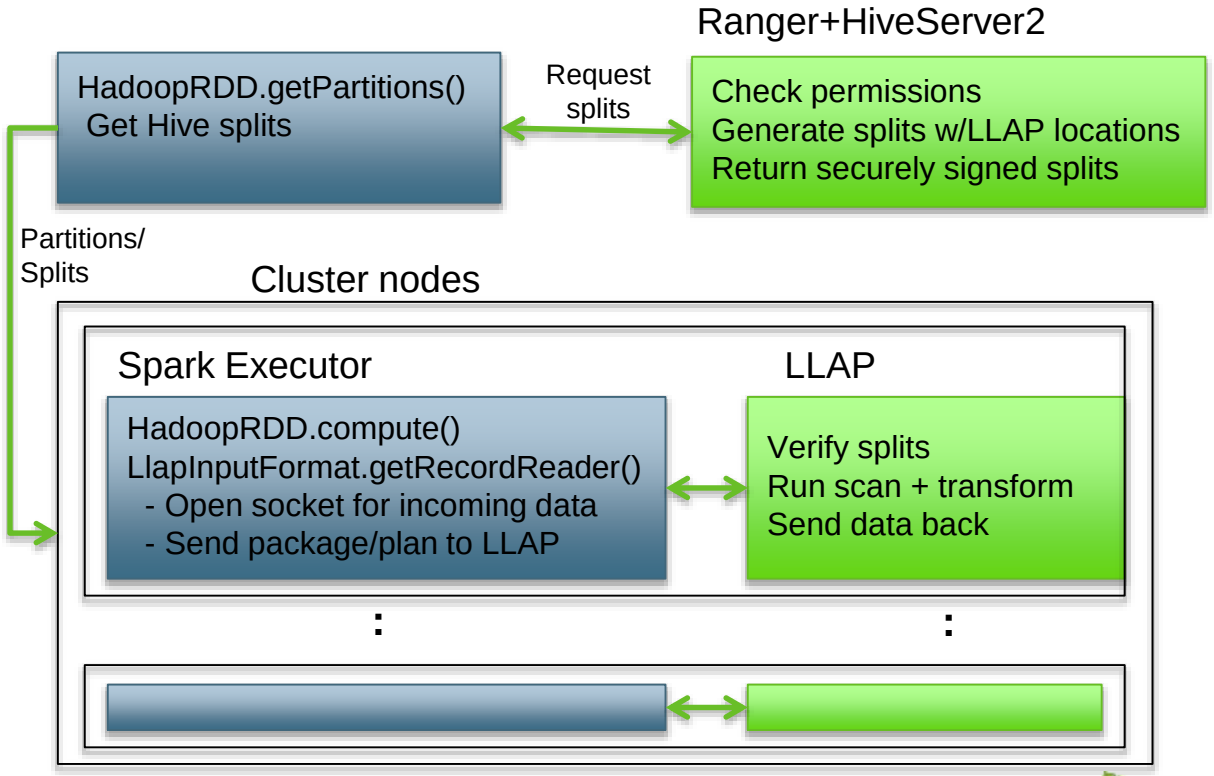

Hortonworks

Monitoring LLAP Queries

Monitoring

- LLAP exposes a UI for monitoring
- Also has jmx endpoint with much more data, logs and jstack endpoints as usual
- Aggregate monitoring UI is work in progress

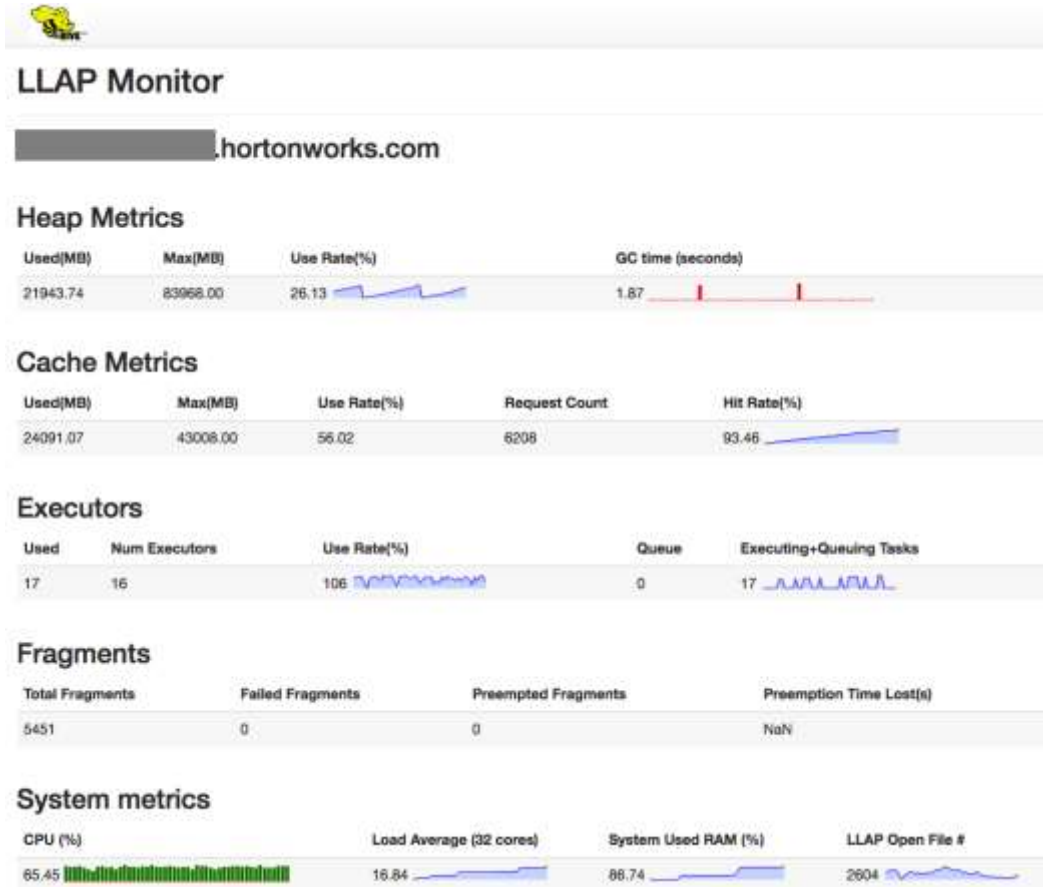

Hortonworks

Watching queries – Tez UI integration

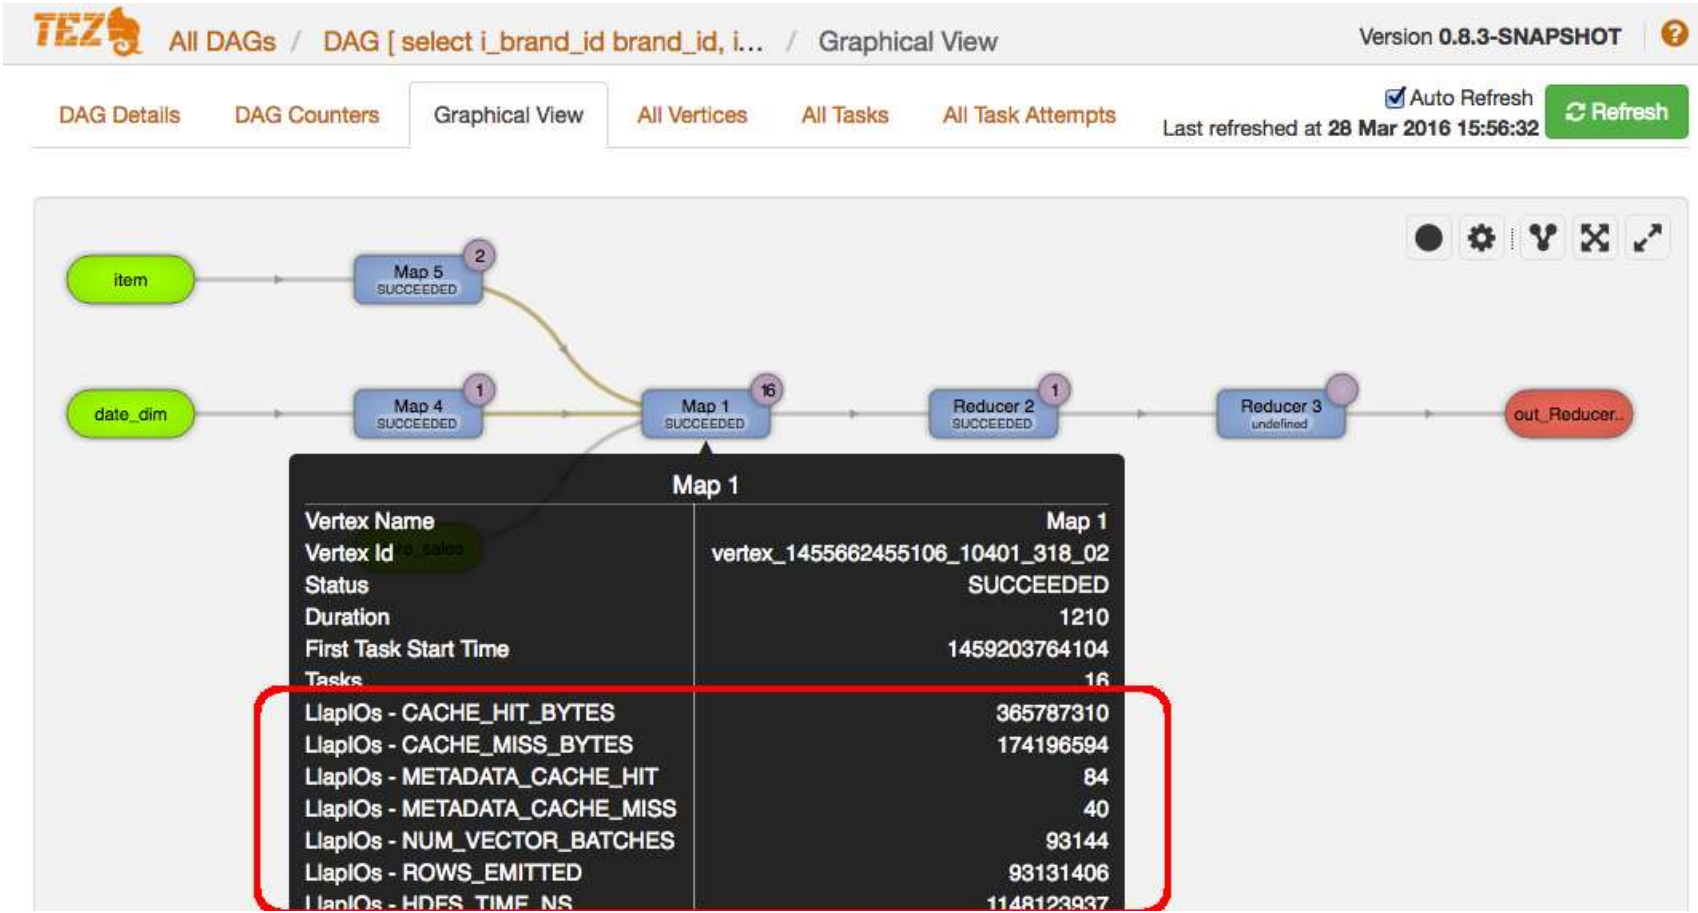

Questions?

Interested? Stop by the Hortonworks booth to learn more

Hortonworks